

# Nueral Network Assi

Pankajan T.

## Libraries used

```
In [ ]: import numpy as np
import pandas as pd
%matplotlib inline
import matplotlib as mpl
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.datasets import cifar10
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.layers import Conv2D, BatchNormalization, Activation, Add, AveragePooling2D
from tensorflow.keras.regularizers import l2
from tensorflow.keras.models import Model
from tensorflow.keras.applications import VGG16
from tensorflow.keras import models
from tensorflow.keras import layers
from tensorflow.keras import optimizers
from tensorflow.keras.applications import Xception
```

## Dataset used

Here we use CIFAR10 of keras dataset that have extensive images of objects and animals containing 10 classes given below.

'airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck'

## Goal

The main objective is to develop the model that can classify images that are each of the above given label. We design the model to be more accurate as possible.

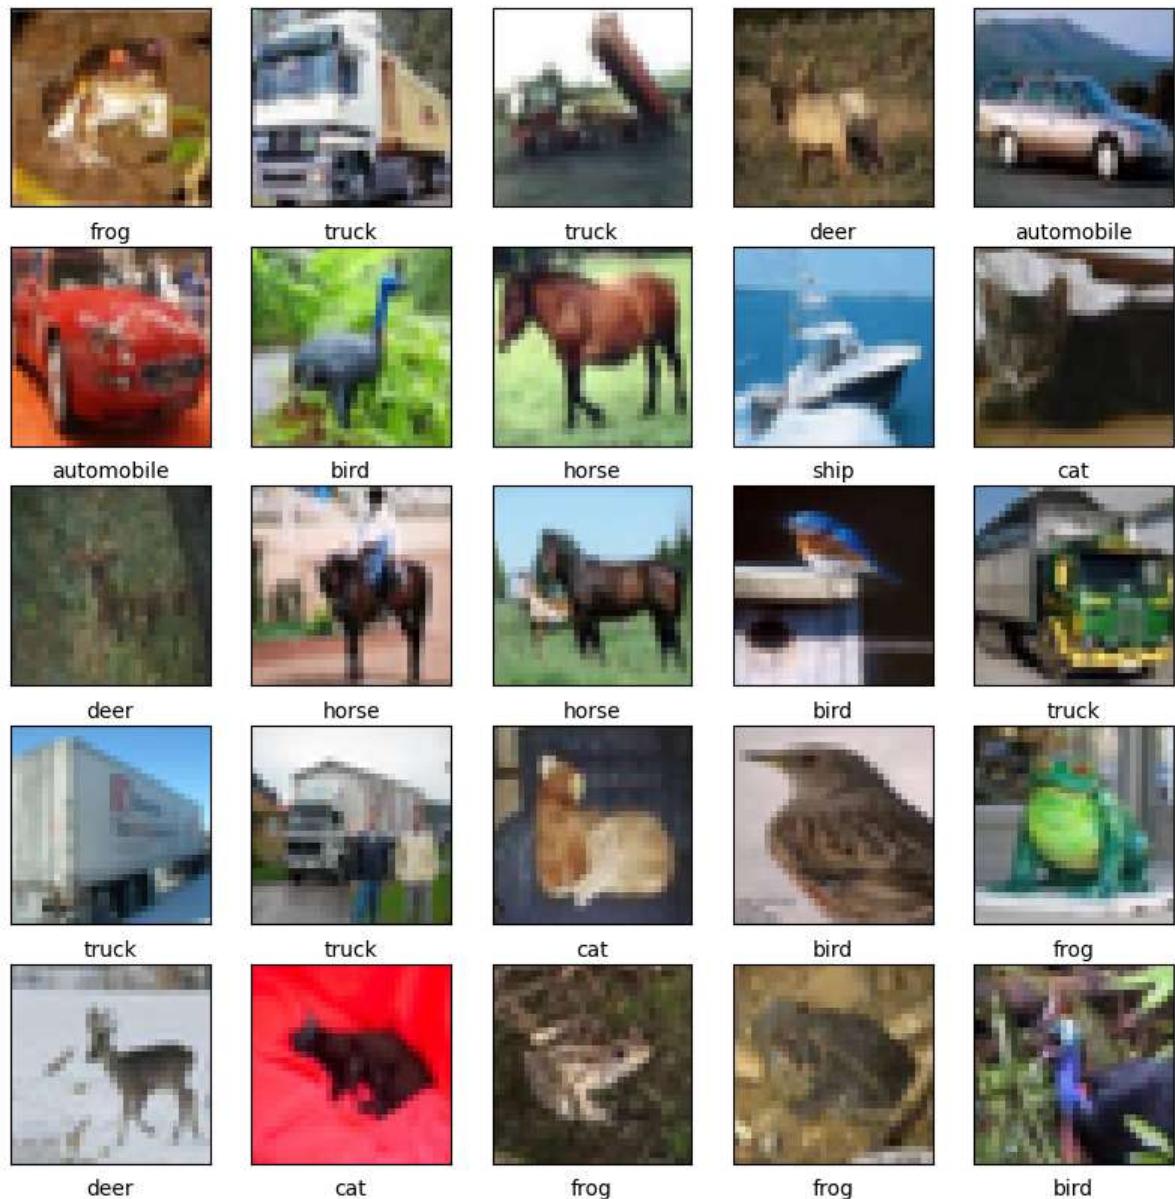
## Data preprocessing

```
In [ ]: (x_train, y_train), (x_test, y_test) = keras.datasets.cifar10.load_data()
        x_train, x_test = x_train / 255.0, x_test / 255.0

In [ ]: class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',
                    'dog', 'frog', 'horse', 'ship', 'truck']

In [ ]: plt.figure(figsize=(10,10))
        for i in range(25):
```

```
plt.subplot(5,5,i+1)
plt.xticks([])
plt.yticks([])
plt.grid(False)
plt.imshow(x_train[i])
plt.xlabel(class_names[y_train[i][0]])
plt.show()
```



## Create the model architecture

```
In [ ]: np.random.seed(42)
tf.random.set_seed(42)
```

### 1. CNN-1

```
In [ ]: model = keras.models.Sequential()
model.add(keras.layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3))
```

```
model.add(keras.layers.MaxPooling2D((2, 2)))
model.add(keras.layers.Conv2D(64, (3, 3), activation='relu'))
model.add(keras.layers.MaxPooling2D((2, 2)))
model.add(keras.layers.Conv2D(64, (3, 3), activation='relu'))
```

In [ ]: `model.summary()`

Model: "sequential\_9"

Layer (type)	Output Shape	Param #
<hr/>		
Layer (type)	Output Shape	Param #
conv2d_11 (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d_7 (MaxPooling 2D)	(None, 15, 15, 32)	0
conv2d_12 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_8 (MaxPooling 2D)	(None, 6, 6, 64)	0
conv2d_13 (Conv2D)	(None, 4, 4, 64)	36928
<hr/>		
Total params: 56,320		
Trainable params: 56,320		
Non-trainable params: 0		

In [ ]: `model.add(keras.layers.Flatten())
model.add(keras.layers.Dense(64, activation='relu'))
model.add(keras.layers.Dense(10))`

In [ ]: `model.summary()`

Model: "sequential\_9"

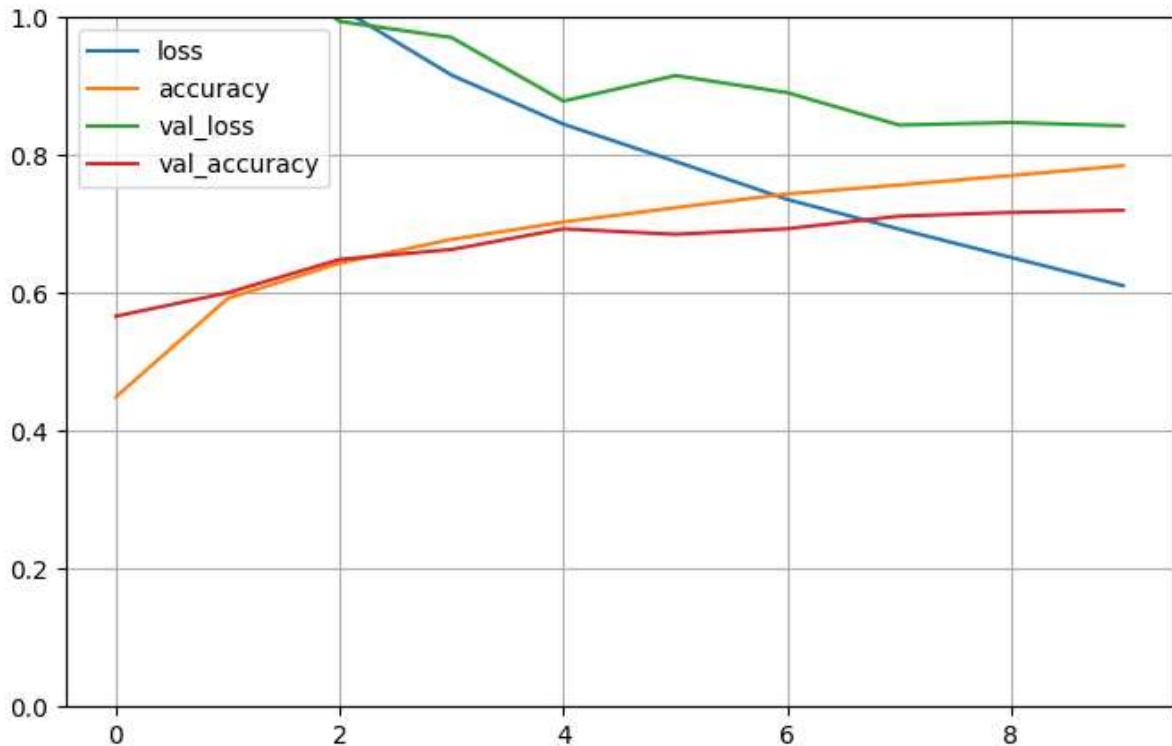
Layer (type)	Output Shape	Param #
<hr/>		
conv2d_11 (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d_7 (MaxPooling 2D)	(None, 15, 15, 32)	0
conv2d_12 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_8 (MaxPooling 2D)	(None, 6, 6, 64)	0
<hr/>		
Layer (type)	Output Shape	Param #
conv2d_11 (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d_7 (MaxPooling 2D)	(None, 15, 15, 32)	0
conv2d_12 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_8 (MaxPooling 2D)	(None, 6, 6, 64)	0
conv2d_13 (Conv2D)	(None, 4, 4, 64)	36928
flatten_6 (Flatten)	(None, 1024)	0
dense_17 (Dense)	(None, 64)	65600
dense_18 (Dense)	(None, 10)	650
<hr/>		
Total params: 122,570		
Trainable params: 122,570		
Non-trainable params: 0		

```
In [ ]: model.compile(optimizer='adam',
                      loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
                      metrics=['accuracy'])
```

```
In [ ]: history = model.fit(x_train, y_train, epochs=10,
                           validation_data=(x_test, y_test))
```

```
Epoch 1/10
1563/1563 [=====] - 64s 39ms/step - loss: 1.5172 - accuracy: 0.4492 - val_loss: 1.2107 - val_accuracy: 0.5664
Epoch 2/10
1563/1563 [=====] - 62s 40ms/step - loss: 1.1524 - accuracy: 0.5923 - val_loss: 1.1139 - val_accuracy: 0.6004
Epoch 3/10
1563/1563 [=====] - 52s 33ms/step - loss: 1.0142 - accuracy: 0.6434 - val_loss: 0.9941 - val_accuracy: 0.6488
Epoch 4/10
1563/1563 [=====] - 72s 46ms/step - loss: 0.9161 - accuracy: 0.6779 - val_loss: 0.9707 - val_accuracy: 0.6632
Epoch 5/10
1563/1563 [=====] - 100s 64ms/step - loss: 0.8449 - accuracy: 0.7031 - val_loss: 0.8785 - val_accuracy: 0.6932
Epoch 6/10
1563/1563 [=====] - 106s 68ms/step - loss: 0.7910 - accuracy: 0.7238 - val_loss: 0.9155 - val_accuracy: 0.6852
Epoch 7/10
1563/1563 [=====] - 119s 76ms/step - loss: 0.7358 - accuracy: 0.7439 - val_loss: 0.8906 - val_accuracy: 0.6933
Epoch 8/10
1563/1563 [=====] - 88s 56ms/step - loss: 0.6928 - accuracy: 0.7568 - val_loss: 0.8437 - val_accuracy: 0.7116
Epoch 9/10
1563/1563 [=====] - 79s 51ms/step - loss: 0.6517 - accuracy: 0.7705 - val_loss: 0.8474 - val_accuracy: 0.7169
Epoch 10/10
1563/1563 [=====] - 78s 50ms/step - loss: 0.6107 - accuracy: 0.7848 - val_loss: 0.8426 - val_accuracy: 0.7200
```

```
In [ ]: pd.DataFrame(history.history).plot(figsize=(8, 5))
plt.grid(True)
plt.gca().set_ylim(0, 1)
plt.show()
```



```
In [ ]: ev = model.evaluate(x_test, y_test)
```

```
313/313 [=====] - 6s 18ms/step - loss: 0.8426 - accuracy: 0.7200
```

```
In [ ]: ev
```

```
Out[ ]: [0.842645525932312, 0.7200000286102295]
```

```
In [ ]: keras.backend.clear_session
```

```
Out[ ]: <function keras.backend.clear_session()>
```

```
In [ ]: del model
```

## 2. CNN with pooling

```
In [ ]: model_a = keras.models.Sequential()
model_a.add(keras.layers.Conv2D(filters = 32, kernel_size = (3, 3), strides=1, padding='same'))
model_a.add(keras.layers.MaxPooling2D((2, 2)))
model_a.add(keras.layers.Flatten())
model_a.add(keras.layers.Dense(300, activation="relu"))
model_a.add(keras.layers.Dense(100, activation="relu"))
model_a.add(keras.layers.Dense(10, activation="softmax"))
```

```
In [ ]: model_a.summary()
```

Model: "sequential\_20"

Layer (type)	Output Shape	Param #
<hr/>		
conv2d_112 (Conv2D)	(None, 30, 30, 32)	896
<hr/>		
conv2d_112 (Conv2D)	(None, 30, 30, 32)	896
<hr/>		
max_pooling2d_11 (MaxPooling2D)	(None, 15, 15, 32)	0
<hr/>		
flatten_15 (Flatten)	(None, 7200)	0
<hr/>		
dense_40 (Dense)	(None, 300)	2160300
<hr/>		
dense_41 (Dense)	(None, 100)	30100
<hr/>		
dense_42 (Dense)	(None, 10)	1010
<hr/>		
Total params: 2,192,306		
Trainable params: 2,192,306		
Non-trainable params: 0		

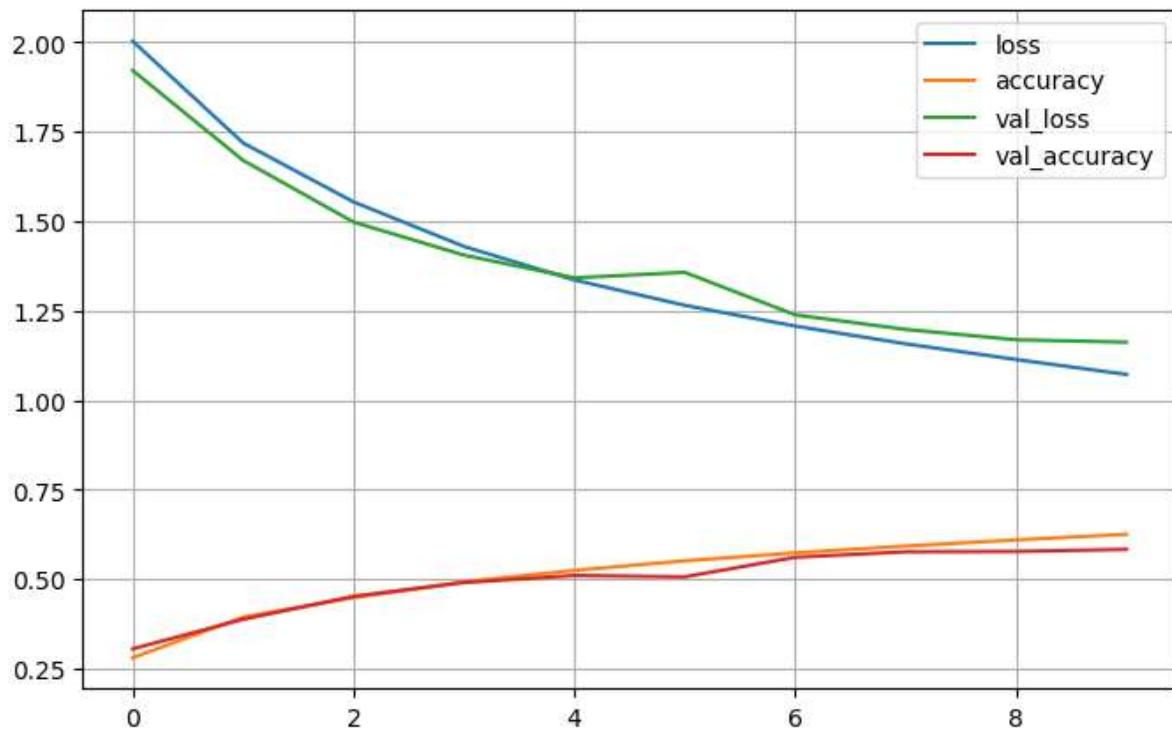
```
In [ ]: model_a.compile(loss="sparse_categorical_crossentropy",
                      optimizer="sgd",
                      metrics=["accuracy"])
```

```
In [ ]: model_history_a = model_a.fit(x_train,y_train, epochs=20,batch_size= 64,
                                      validation_data=(x_test, y_test))
```

Epoch 1/20  
782/782 [=====] - 42s 48ms/step - loss: 0.8253 - accuracy: 0.7142 - val\_loss: 1.1542 - val\_accuracy: 0.6026  
Epoch 2/20  
782/782 [=====] - 37s 47ms/step - loss: 0.7874 - accuracy: 0.7271 - val\_loss: 1.1062 - val\_accuracy: 0.6202  
Epoch 3/20  
782/782 [=====] - 36s 45ms/step - loss: 0.7520 - accuracy: 0.7380 - val\_loss: 1.1140 - val\_accuracy: 0.6140  
Epoch 4/20  
782/782 [=====] - 32s 41ms/step - loss: 0.7165 - accuracy: 0.7542 - val\_loss: 1.0424 - val\_accuracy: 0.6421  
Epoch 5/20  
782/782 [=====] - 28s 36ms/step - loss: 0.6830 - accuracy: 0.7620 - val\_loss: 1.2674 - val\_accuracy: 0.5852  
Epoch 6/20  
782/782 [=====] - 30s 38ms/step - loss: 0.6450 - accuracy: 0.7790 - val\_loss: 1.2177 - val\_accuracy: 0.6073  
Epoch 7/20  
782/782 [=====] - 39s 49ms/step - loss: 0.6115 - accuracy: 0.7902 - val\_loss: 1.0817 - val\_accuracy: 0.6402  
Epoch 8/20  
782/782 [=====] - 35s 45ms/step - loss: 0.5781 - accuracy: 0.8036 - val\_loss: 1.0546 - val\_accuracy: 0.6547  
Epoch 9/20  
782/782 [=====] - 37s 48ms/step - loss: 0.5417 - accuracy: 0.8149 - val\_loss: 1.1382 - val\_accuracy: 0.6339  
Epoch 10/20  
782/782 [=====] - 43s 55ms/step - loss: 0.5070 - accuracy: 0.8286 - val\_loss: 1.2070 - val\_accuracy: 0.6118  
Epoch 11/20  
782/782 [=====] - 36s 46ms/step - loss: 0.4713 - accuracy: 0.8427 - val\_loss: 1.2342 - val\_accuracy: 0.6220  
Epoch 12/20  
782/782 [=====] - 34s 43ms/step - loss: 0.4375 - accuracy: 0.8546 - val\_loss: 1.3746 - val\_accuracy: 0.5960  
Epoch 13/20  
782/782 [=====] - 36s 46ms/step - loss: 0.4031 - accuracy: 0.8684 - val\_loss: 1.1688 - val\_accuracy: 0.6453  
Epoch 14/20  
782/782 [=====] - 33s 42ms/step - loss: 0.3661 - accuracy: 0.8825 - val\_loss: 1.4833 - val\_accuracy: 0.5961  
Epoch 15/20  
782/782 [=====] - 33s 42ms/step - loss: 0.3368 - accuracy: 0.8938 - val\_loss: 1.1873 - val\_accuracy: 0.6524  
Epoch 16/20  
782/782 [=====] - 31s 40ms/step - loss: 0.3049 - accuracy: 0.9044 - val\_loss: 2.8354 - val\_accuracy: 0.4862  
Epoch 17/20  
782/782 [=====] - 28s 36ms/step - loss: 0.2794 - accuracy: 0.9147 - val\_loss: 1.4298 - val\_accuracy: 0.6235  
Epoch 18/20  
782/782 [=====] - 32s 41ms/step - loss: 0.2463 - accuracy: 0.9267 - val\_loss: 1.2798 - val\_accuracy: 0.6474  
Epoch 19/20  
782/782 [=====] - 35s 45ms/step - loss: 0.2186 - accuracy:

```
y: 0.9370 - val_loss: 1.4434 - val_accuracy: 0.6375
Epoch 20/20
782/782 [=====] - 34s 43ms/step - loss: 0.1933 - accurac
y: 0.9468 - val_loss: 1.4510 - val_accuracy: 0.6312
```

```
In [ ]: pd.DataFrame(model_history_a.history).plot(figsize=(8, 5))
plt.grid(True)
plt.show()
```



### 3. VGG16

```
In [ ]: conv_base = VGG16(weights='imagenet',
                         include_top=False,
                         input_shape=(32, 32, 3))
```

```
In [ ]: conv_base.summary()
```

Model: "vgg16"

Layer (type)	Output Shape	Param #
<hr/>		
Layer (type)	Output Shape	Param #
input_5 (InputLayer)	[None, 32, 32, 3]	0
block1_conv1 (Conv2D)	(None, 32, 32, 64)	1792
block1_conv2 (Conv2D)	(None, 32, 32, 64)	36928
block1_pool (MaxPooling2D)	(None, 16, 16, 64)	0
block2_conv1 (Conv2D)	(None, 16, 16, 128)	73856
block2_conv2 (Conv2D)	(None, 16, 16, 128)	147584
block2_pool (MaxPooling2D)	(None, 8, 8, 128)	0
block3_conv1 (Conv2D)	(None, 8, 8, 256)	295168
block3_conv2 (Conv2D)	(None, 8, 8, 256)	590080
block3_conv3 (Conv2D)	(None, 8, 8, 256)	590080
block3_pool (MaxPooling2D)	(None, 4, 4, 256)	0
block4_conv1 (Conv2D)	(None, 4, 4, 512)	1180160
block4_conv2 (Conv2D)	(None, 4, 4, 512)	2359808
block4_conv3 (Conv2D)	(None, 4, 4, 512)	2359808
block4_pool (MaxPooling2D)	(None, 2, 2, 512)	0
block5_conv1 (Conv2D)	(None, 2, 2, 512)	2359808
block5_conv2 (Conv2D)	(None, 2, 2, 512)	2359808
block5_conv3 (Conv2D)	(None, 2, 2, 512)	2359808
block5_pool (MaxPooling2D)	(None, 1, 1, 512)	0
<hr/>		
Total params: 14,714,688		
Trainable params: 14,714,688		
Non-trainable params: 0		

---

```
In [ ]: model = models.Sequential()
model.add(conv_base)
model.add(layers.Flatten())
model.add(layers.Dense(256, activation='relu'))
model.add(layers.Dense(10, activation='sigmoid'))
```

```
In [ ]: model.summary()
```

Model: "sequential\_12"

Layer (type)	Output Shape	Param #
<hr/>		
vgg16 (Functional)	(None, 1, 1, 512)	14714688
flatten_12 (Flatten)	(None, 512)	0
dense_27 (Dense)	(None, 256)	131328
<hr/>		
Layer (type)	Output Shape	Param #
<hr/>		
vgg16 (Functional)	(None, 1, 1, 512)	14714688
flatten_12 (Flatten)	(None, 512)	0
dense_27 (Dense)	(None, 256)	131328
dense_28 (Dense)	(None, 10)	2570
<hr/>		
Total params: 14,848,586		
Trainable params: 133,898		
Non-trainable params: 14,714,688		

```
In [ ]: conv_base.trainable = False
```

```
In [ ]: model.compile(loss='binary_crossentropy',
                      optimizer=optimizers.RMSprop(lr=2e-5),
                      metrics=['acc'])
```

WARNING:absl:`lr` is deprecated, please use `learning\_rate` instead, or use the legacy optimizer, e.g.,tf.keras.optimizers.RMSprop.

```
In [ ]: checkpoint_cb = keras.callbacks.ModelCheckpoint("CNN_Project_Model-{epoch:02d}.h5")
```

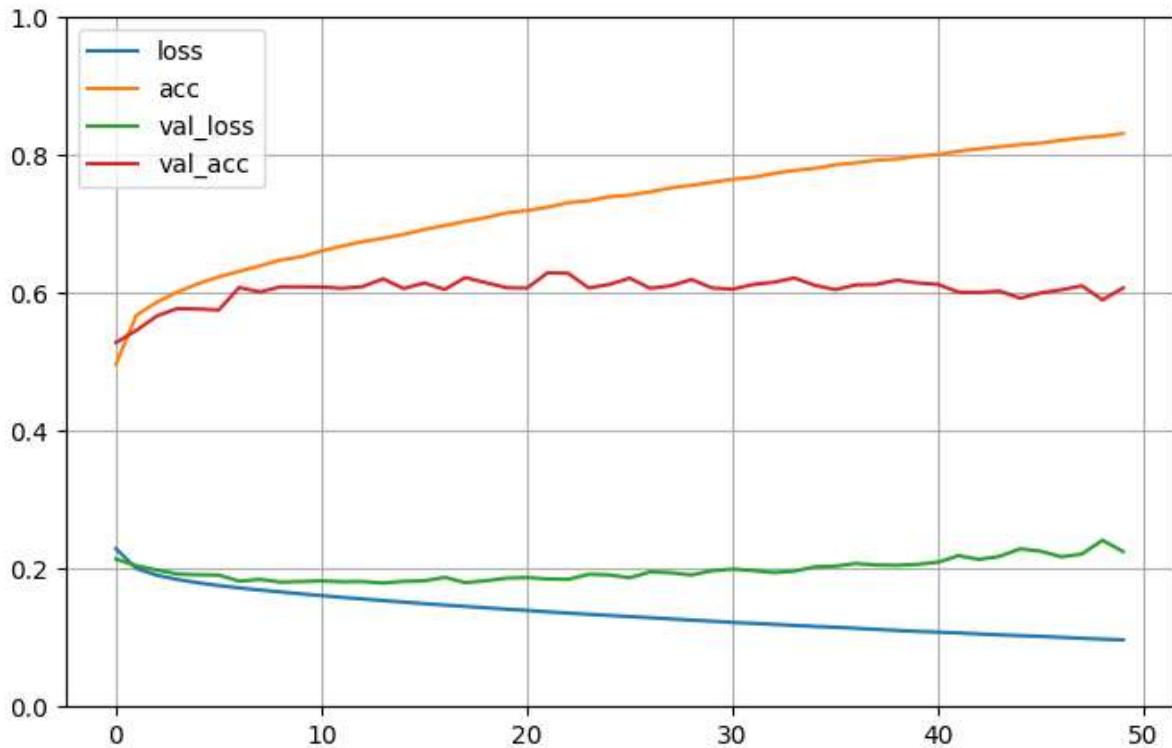
```
In [ ]: history = model.fit(x_train,y_train, epochs=50,batch_size= 64,
                           validation_data=(x_test, y_test),
                           validation_steps=50,
                           callbacks=[checkpoint_cb])
```

```
Epoch 1/50
782/782 [=====] - 154s 192ms/step - loss: 0.2292 - acc: 0.4963 - val_loss: 0.2141 - val_acc: 0.5281
Epoch 2/50
782/782 [=====] - 171s 219ms/step - loss: 0.2003 - acc: 0.5681 - val_loss: 0.2043 - val_acc: 0.5459
Epoch 3/50
782/782 [=====] - 193s 246ms/step - loss: 0.1906 - acc: 0.5869 - val_loss: 0.1982 - val_acc: 0.5669
Epoch 4/50
782/782 [=====] - 196s 250ms/step - loss: 0.1844 - acc: 0.6017 - val_loss: 0.1923 - val_acc: 0.5775
Epoch 5/50
782/782 [=====] - 170s 218ms/step - loss: 0.1797 - acc: 0.6136 - val_loss: 0.1913 - val_acc: 0.5769
Epoch 6/50
782/782 [=====] - 205s 262ms/step - loss: 0.1757 - acc: 0.6236 - val_loss: 0.1908 - val_acc: 0.5753
Epoch 7/50
782/782 [=====] - 110s 141ms/step - loss: 0.1722 - acc: 0.6315 - val_loss: 0.1820 - val_acc: 0.6081
Epoch 8/50
782/782 [=====] - 112s 143ms/step - loss: 0.1691 - acc: 0.6393 - val_loss: 0.1846 - val_acc: 0.6016
Epoch 9/50
782/782 [=====] - 111s 142ms/step - loss: 0.1664 - acc: 0.6481 - val_loss: 0.1807 - val_acc: 0.6091
Epoch 10/50
782/782 [=====] - 112s 143ms/step - loss: 0.1636 - acc: 0.6526 - val_loss: 0.1814 - val_acc: 0.6091
Epoch 11/50
782/782 [=====] - 113s 144ms/step - loss: 0.1610 - acc: 0.6610 - val_loss: 0.1823 - val_acc: 0.6087
Epoch 12/50
782/782 [=====] - 111s 143ms/step - loss: 0.1585 - acc: 0.6681 - val_loss: 0.1811 - val_acc: 0.6069
Epoch 13/50
782/782 [=====] - 111s 142ms/step - loss: 0.1562 - acc: 0.6746 - val_loss: 0.1815 - val_acc: 0.6094
Epoch 14/50
782/782 [=====] - 111s 142ms/step - loss: 0.1539 - acc: 0.6795 - val_loss: 0.1792 - val_acc: 0.6206
Epoch 15/50
782/782 [=====] - 112s 143ms/step - loss: 0.1516 - acc: 0.6851 - val_loss: 0.1818 - val_acc: 0.6066
Epoch 16/50
782/782 [=====] - 111s 142ms/step - loss: 0.1494 - acc: 0.6922 - val_loss: 0.1824 - val_acc: 0.6147
Epoch 17/50
782/782 [=====] - 111s 142ms/step - loss: 0.1473 - acc: 0.6980 - val_loss: 0.1875 - val_acc: 0.6053
Epoch 18/50
782/782 [=====] - 111s 142ms/step - loss: 0.1453 - acc: 0.7040 - val_loss: 0.1798 - val_acc: 0.6225
Epoch 19/50
782/782 [=====] - 111s 142ms/step - loss: 0.1434 - acc:
```

```
0.7095 - val_loss: 0.1826 - val_acc: 0.6150
Epoch 20/50
782/782 [=====] - 111s 142ms/step - loss: 0.1413 - acc:
0.7163 - val_loss: 0.1863 - val_acc: 0.6078
Epoch 21/50
782/782 [=====] - 112s 143ms/step - loss: 0.1396 - acc:
0.7197 - val_loss: 0.1872 - val_acc: 0.6072
Epoch 22/50
782/782 [=====] - 112s 143ms/step - loss: 0.1375 - acc:
0.7245 - val_loss: 0.1852 - val_acc: 0.6294
Epoch 23/50
782/782 [=====] - 112s 143ms/step - loss: 0.1357 - acc:
0.7311 - val_loss: 0.1846 - val_acc: 0.6288
Epoch 24/50
782/782 [=====] - 112s 143ms/step - loss: 0.1339 - acc:
0.7339 - val_loss: 0.1920 - val_acc: 0.6075
Epoch 25/50
782/782 [=====] - 112s 143ms/step - loss: 0.1323 - acc:
0.7398 - val_loss: 0.1910 - val_acc: 0.6125
Epoch 26/50
782/782 [=====] - 111s 142ms/step - loss: 0.1306 - acc:
0.7421 - val_loss: 0.1870 - val_acc: 0.6216
Epoch 27/50
782/782 [=====] - 116s 148ms/step - loss: 0.1290 - acc:
0.7468 - val_loss: 0.1955 - val_acc: 0.6069
Epoch 28/50
782/782 [=====] - 119s 152ms/step - loss: 0.1272 - acc:
0.7525 - val_loss: 0.1943 - val_acc: 0.6106
Epoch 29/50
782/782 [=====] - 119s 152ms/step - loss: 0.1255 - acc:
0.7563 - val_loss: 0.1910 - val_acc: 0.6197
Epoch 30/50
782/782 [=====] - 119s 152ms/step - loss: 0.1239 - acc:
0.7606 - val_loss: 0.1970 - val_acc: 0.6075
Epoch 31/50
782/782 [=====] - 119s 152ms/step - loss: 0.1222 - acc:
0.7651 - val_loss: 0.1995 - val_acc: 0.6056
Epoch 32/50
782/782 [=====] - 119s 152ms/step - loss: 0.1209 - acc:
0.7679 - val_loss: 0.1973 - val_acc: 0.6125
Epoch 33/50
782/782 [=====] - 119s 152ms/step - loss: 0.1194 - acc:
0.7735 - val_loss: 0.1945 - val_acc: 0.6156
Epoch 34/50
782/782 [=====] - 131s 168ms/step - loss: 0.1178 - acc:
0.7779 - val_loss: 0.1965 - val_acc: 0.6219
Epoch 35/50
782/782 [=====] - 126s 161ms/step - loss: 0.1164 - acc:
0.7808 - val_loss: 0.2027 - val_acc: 0.6112
Epoch 36/50
782/782 [=====] - 152s 195ms/step - loss: 0.1151 - acc:
0.7861 - val_loss: 0.2037 - val_acc: 0.6050
Epoch 37/50
782/782 [=====] - 135s 173ms/step - loss: 0.1136 - acc:
0.7889 - val_loss: 0.2077 - val_acc: 0.6119
Epoch 38/50
```

```
782/782 [=====] - 125s 160ms/step - loss: 0.1121 - acc: 0.7927 - val_loss: 0.2055 - val_acc: 0.6125
Epoch 39/50
782/782 [=====] - 137s 176ms/step - loss: 0.1105 - acc: 0.7944 - val_loss: 0.2051 - val_acc: 0.6187
Epoch 40/50
782/782 [=====] - 127s 162ms/step - loss: 0.1091 - acc: 0.7987 - val_loss: 0.2064 - val_acc: 0.6147
Epoch 41/50
782/782 [=====] - 129s 166ms/step - loss: 0.1081 - acc: 0.8010 - val_loss: 0.2096 - val_acc: 0.6125
Epoch 42/50
782/782 [=====] - 143s 183ms/step - loss: 0.1069 - acc: 0.8059 - val_loss: 0.2190 - val_acc: 0.6012
Epoch 43/50
782/782 [=====] - 154s 197ms/step - loss: 0.1054 - acc: 0.8093 - val_loss: 0.2136 - val_acc: 0.6009
Epoch 44/50
782/782 [=====] - 156s 200ms/step - loss: 0.1040 - acc: 0.8122 - val_loss: 0.2180 - val_acc: 0.6028
Epoch 45/50
782/782 [=====] - 152s 194ms/step - loss: 0.1029 - acc: 0.8154 - val_loss: 0.2290 - val_acc: 0.5925
Epoch 46/50
782/782 [=====] - 135s 173ms/step - loss: 0.1019 - acc: 0.8176 - val_loss: 0.2254 - val_acc: 0.6003
Epoch 47/50
782/782 [=====] - 131s 167ms/step - loss: 0.1005 - acc: 0.8217 - val_loss: 0.2173 - val_acc: 0.6047
Epoch 48/50
782/782 [=====] - 141s 180ms/step - loss: 0.0992 - acc: 0.8251 - val_loss: 0.2215 - val_acc: 0.6106
Epoch 49/50
782/782 [=====] - 136s 174ms/step - loss: 0.0979 - acc: 0.8275 - val_loss: 0.2413 - val_acc: 0.5900
Epoch 50/50
782/782 [=====] - 138s 177ms/step - loss: 0.0969 - acc: 0.8314 - val_loss: 0.2250 - val_acc: 0.6075
```

```
In [ ]: pd.DataFrame(history.history).plot(figsize=(8, 5))
plt.grid(True)
plt.gca().set_ylim(0, 1)
plt.show()
```



```
In [ ]: hist_df = pd.DataFrame(history.history)
```

```
In [ ]: results = model.evaluate(x_test, y_test, batch_size=128)
```

```
79/79 [=====] - 21s 265ms/step - loss: 0.2277 - acc: 0.6117
```

## 4. Xception

```
In [ ]: base_model = Xception(include_top=False, weights='imagenet', input_shape=(224,224,3))
```

```
Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/xception/xception_weights_tf_dim_ordering_tf_kernels_notop.h5
83683744/83683744 [=====] - 74s 1us/step
```

```
In [ ]: base_model.trainable = False
```

```
In [ ]: from tensorflow.keras.layers import Input, Flatten, Dense, BatchNormalization, Activation
data_augmentation = Sequential([
    RandomFlip("horizontal"),
    RandomRotation(0.1),
    RandomZoom(0.1)
])
```

```
In [ ]: inputs = tf.keras.Input(shape=(32, 32, 3))
x = tf.keras.layers.Lambda(lambda image: tf.image.resize(image, (224,224)))(inputs)
x = tf.keras.applications.xception.preprocess_input(x)
x = base_model(x, training=False)
x = tf.keras.layers.GlobalAveragePooling2D()(x)
x = tf.keras.layers.Dropout(0.3)(x)
```

```
outputs = tf.keras.layers.Dense(10, activation='softmax'))(x)
model = tf.keras.Model(inputs, outputs)
```

In [ ]: `model.summary()`

Model: "model\_4"

Layer (type)	Output Shape	Param #
Layer (type)	Output Shape	Param #
<hr/>		
input_17 (InputLayer)	[None, 32, 32, 3]	0
lambda_1 (Lambda)	(None, 224, 224, 3)	0
tf.math.truediv_1 (TFOpLambda)	(None, 224, 224, 3)	0
tf.math.subtract_1 (TFOpLambda)	(None, 224, 224, 3)	0
xception (Functional)	(None, 7, 7, 2048)	20861480
global_average_pooling2d_1 (GlobalAveragePooling2D)	(None, 2048)	0
dropout_1 (Dropout)	(None, 2048)	0
dense_36 (Dense)	(None, 10)	20490
<hr/>		
Total params:	20,881,970	
Trainable params:	20,490	
Non-trainable params:	20,861,480	

In [ ]: `model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])`

In [ ]: `epochs = 10`  
`history = model.fit(x_train, y_train, validation_data=(x_test, y_test), epochs=epochs)`

```
Epoch 1/10
1563/1563 [=====] - 3975s 3s/step - loss: 2.2943 - accuracy: 0.1226 - val_loss: 2.2764 - val_accuracy: 0.1372
Epoch 2/10
1563/1563 [=====] - 3193s 2s/step - loss: 2.2697 - accuracy: 0.1534 - val_loss: 2.2558 - val_accuracy: 0.2203
Epoch 3/10
717/1563 [=====>.....] - ETA: 9:23:26 - loss: 2.2555 - accuracy: 0.1712
```

```

-----
KeyboardInterrupt                                     Traceback (most recent call last)

Cell In[151], line 2
    1 epochs = 10
----> 2 history = model.fit(x_train, y_train, validation_data=(x_test, y_test), ep
ochs=epochs, verbose=1)

File e:\programes\python\lib\site-packages\keras\utils\traceback_utils.py:65, in f
ilter_traceback.<locals>.error_handler(*args, **kwargs)
    63 filtered_tb = None
    64 try:
--> 65     return fn(*args, **kwargs)
    66 except Exception as e:
    67     filtered_tb = _process_traceback_frames(e.__traceback__)

File e:\programes\python\lib\site-packages\keras\engine\training.py:1650, in Mode
l.fit(self, x, y, batch_size, epochs, verbose, callbacks, validation_split, valida
tion_data, shuffle, class_weight, sample_weight, initial_epoch, steps_per_epoch, v
alidation_steps, validation_batch_size, validation_freq, max_queue_size, workers,
use_multiprocessing)
    1642 with tf.profiler.experimental.Trace(
    1643     "train",
    1644     epoch_num=epoch,
    (...),
    1647     _r=1,
    1648 ):
    1649     callbacks.on_train_batch_begin(step)
-> 1650     tmp_logs = self.train_function(iterator)
    1651     if data_handler.should_sync:
    1652         context.async_wait()

File e:\programes\python\lib\site-packages\tensorflow\python\util\traceback_utils.
py:150, in filter_traceback.<locals>.error_handler(*args, **kwargs)
    148 filtered_tb = None
    149 try:
--> 150     return fn(*args, **kwargs)
    151 except Exception as e:
    152     filtered_tb = _process_traceback_frames(e.__traceback__)

File e:\programes\python\lib\site-packages\tensorflow\python\eager\polymorphic_fun
ction\polymorphic_function.py:880, in Function.__call__(self, *args, **kwds)
    877 compiler = "xla" if self._jit_compile else "nonXla"
    879 with OptionalXlaContext(self._jit_compile):
--> 880     result = self._call(*args, **kwds)
    882 new_tracing_count = self.experimental_get_tracing_count()
    883 without_tracing = (tracing_count == new_tracing_count)

File e:\programes\python\lib\site-packages\tensorflow\python\eager\polymorphic_fun
ction\polymorphic_function.py:912, in Function._call(self, *args, **kwds)
    909     self._lock.release()
    910     # In this case we have created variables on the first call, so we run th
e
    911     # defunned version which is guaranteed to never create variables.
--> 912     return self._no_variable_creation_fn(*args, **kwds) # pylint: disable=n
ot-callable
    913 elif self._variable_creation_fn is not None:

```

```

914     # Release the lock early so that multiple threads can perform the call
915     # in parallel.
916     self._lock.release()

File e:\programes\python\lib\site-packages\tensorflow\python\eager\polymorphic_function\tracing_compiler.py:134, in TracingCompiler.__call__(self, *args, **kwargs)
 131     with self._lock:
 132         (concrete_function,
 133          filtered_flat_args) = self._maybe_define_function(args, kwargs)
--> 134     return concrete_function._call_flat(
 135             filtered_flat_args, captured_inputs=concrete_function.captured_inputs)

File e:\programes\python\lib\site-packages\tensorflow\python\eager\polymorphic_function\monomorphic_function.py:1745, in ConcreteFunction._call_flat(self, args, captured_inputs, cancellation_manager)
 1741     possible_gradient_type = gradients_util.PossibleTapeGradientTypes(args)
 1742     if (possible_gradient_type == gradients_util.POSSIBLE_GRADIENT_TYPES_NONE
 1743         and executing_eagerly):
 1744         # No tape is watching; skip to running the function.
-> 1745     return self._build_call_outputs(self._inference_function.call(
 1746         ctx, args, cancellation_manager=cancellation_manager))
 1747     forward_backward = self._select_forward_and_backward_functions(
 1748         args,
 1749         possible_gradient_type,
 1750         executing_eagerly)
 1751     forward_function, args_with_tangents = forward_backward.forward()

File e:\programes\python\lib\site-packages\tensorflow\python\eager\polymorphic_function\monomorphic_function.py:378, in _EagerDefinedFunction.call(self, ctx, args, cancellation_manager)
 376     with _InterpolateFunctionError(self):
 377         if cancellation_manager is None:
--> 378             outputs = execute.execute(
 379                 str(self.signature.name),
 380                 num_outputs=self._num_outputs,
 381                 inputs=args,
 382                 attrs=attrs,
 383                 ctx=ctx)
 384     else:
 385         outputs = execute.execute_with_cancellation(
 386             str(self.signature.name),
 387             num_outputs=self._num_outputs,
(...).
 390             ctx=ctx,
 391             cancellation_manager=cancellation_manager)

File e:\programes\python\lib\site-packages\tensorflow\python\eager\execute.py:52,
in quick_execute(op_name, num_outputs, inputs, attrs, ctx, name)
  50     try:
  51         ctx.ensure_initialized()
--> 52     tensors = pywrap_tfe.TFE_Py_Execute(ctx._handle, device_name, op_name,
  53                                         inputs, attrs, num_outputs)
  54     except core._NotOkStatusException as e:
  55         if name is not None:

```

**KeyboardInterrupt:**

!! Due to running near 10 hours and not yet reaching a promising accuracy level the model have to be stopped mid process

```
In [ ]: results = model.evaluate(x_test, y_test, batch_size=128)
79/79 [=====] - 618s 8s/step - loss: 2.2427 - accuracy: 0.2385
In [ ]: results
Out[ ]: [2.2427217960357666, 0.23849999904632568]
```

## Conclusion

Here the Normal Custom CNN that is implemented without the pooling have the best accuracy of 0.72 and CNN with pooling and VGG16 worked fine with the dataset but the Xception model is too slow, takes very long time and its accuracy is not improving.

## Insights

Here the Dataset works well with CNN models but Xception model is not recommended here. Here the CNN had not much of flaw VGG16 already hit saturation so not further training wont improve the situation.

## Sugestion

Dataset could be expanded to hold more labels and datas. Here VGG16 and Xception models doesn't need further training others could improve. Other complex models can be implemented here.